


Predicting Post-Earnings Announcement Drift

...

Example: Apple Announcing Earning 2010/07/20



Overview

U.S. stocks have been shown to earn higher returns during earnings announcement months than during non-announcement months.

The magnitude of this earnings announcement premium has been estimated by Frazzini and Lamont (2007) to be over 7% per year. While a number of potential explanations for the premium have been put forward, uncertainty remains over the reasons for its existence.

There are two goals of the projects.

- The first is to investigate premium extends over the companies traded in the three major exchanges in the U.S., thereby providing out-of-sample evidence of its existence.
- The second is to exploit observed variations in the direction of the premium in order to gain insights into the factors driving this difference.

Data

Selection and formation of data;
Variables and intuitions

$$\begin{aligned} Ret_{it} = & \alpha_t + \beta_{1t} ifpositiveDrift_{it} + \beta_{2t} Past30Vol_{it} \\ & + \beta_{3t} Past10Vol_{it} + \beta_{4t} PastMktVol_{it} + \beta_{5t} Ret5dbefore_{it} \\ & + \beta_{6t} Ret1dbefore_{it} + \beta_{7t} past200dSK_{it} \\ & + \beta_{8t} past200dKur_{it} + \beta_{9t} PR_{it} + \sum_{j=1}^3 \gamma_{jt} Quarter_{jt} \end{aligned}$$

Data

CRSP Daily Stock

- Period of 2010 - 2018
- Return $\pm 5, 10, 30$ day
- Vol -10 day
- Lag_Returns + 5 day
- Lag_Market Cap
- - 200 day Skewness

Compustat Quarterly

- Earnings report date
- Fiscal Quarter
- Change in Cash
- Shares Repurchase

CRSP Daily Market

- Past 3 Month Market Volatility

Variables

$$\begin{aligned}
 Ret_{it} = & \alpha_t + \beta_{1t}IfPositiveDrift_{it} + \beta_{2t}Past30Vol_{it} \\
 & + \beta_{3t}Past10Vol_{it} + \beta_{4t}PastMktVol_{it} + \beta_{5t}Ret5dbefore_{it} \\
 & + \beta_{6t}Ret1dbefore_{it} + \beta_{7t}past200dSK_{it} \\
 & + \beta_{8t}past200dKur_{it} + \beta_{9t}PR_{it} + \sum_{j=1}^3 \gamma_{jt}Quarter_{jt}
 \end{aligned}$$

IfPositiveDrift, Quarter

- IfPositiveDrift=1, if lag(Ret5dayAfter)>0; IfPositiveDrift=0, if lag(Ret5dayAfter)<0
- Quarter_j: 3 dummies to count the seasonality effect.

PastVol

- PastVol: use past 30 and past 10 days' volatilities as inputs.
- PastMktVol: past quarter's market volatility before announcement day.

PR

- From Ball and Kothari(1991), PR measures investor sophistication.
 $PR_i = [\log(\max MV) - \log(MV_i)] / [\log(\max MV) - \log(\min MV)]$, max MV=\$99 billion, min MV=\$1 million, and MV_i =market value of equity at the beginning of firm-quarter q for firm i , in millions.

past200dSK & past200dKur

- From McNichols(1988), earnings reports causes more 'extreme bad news' to be reflected in stock prices.
- Use past 200 days stock returns before announcement days to calculate skewness and kurtosis.

Models

Panel Regression

R2 score is 3.6%, but this is on individual stocks, not the portfolio

```
Call:
  felm(formula = signs ~ x | fyearq + gvkey | 0 | fyearq + gvkey,      data = w1)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-1.1897 -0.4051  0.0818  0.3753  1.1421
```

```
Coefficients:
                Estimate Cluster s.e. t value Pr(>|t|)
xfqtr          1.179e-02   3.341e-02   0.353   0.7241
xRet5DayBefore  1.327e-01   7.806e-02   1.701   0.0890 .
xPast10dayVol    7.098e-01   4.306e-01   1.648   0.0993 .
xifPositiveDrift_Lag  2.988e-03   3.713e-03   0.805   0.4209
xPast30dayVol   -1.515e+00   6.671e-01  -2.271   0.0232 *
xPastMktVol3Month -3.090e+00   2.195e+00  -1.408   0.1593
xLaggedME       -2.315e-09   1.055e-09  -2.195   0.0282 *
xChangeInCashInvestment -1.064e-07   2.652e-06  -0.040   0.9680
xSharesRepurchased  1.684e-04   7.793e-04   0.216   0.8289
xLTDebtDue1Year   1.577e-06   5.063e-06   0.311   0.7555
xAmortization     2.071e-05   1.082e-04   0.191   0.8483
xRet1DayBefore   -1.384e-01   1.066e-01  -1.298   0.1945
xPast200SK       1.823e-03   2.346e-03   0.777   0.4371
xPast200Kur      -5.138e-04   2.576e-04  -1.995   0.0461 *
xPR              1.452e+00   2.238e-01   6.485   9.1e-11 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

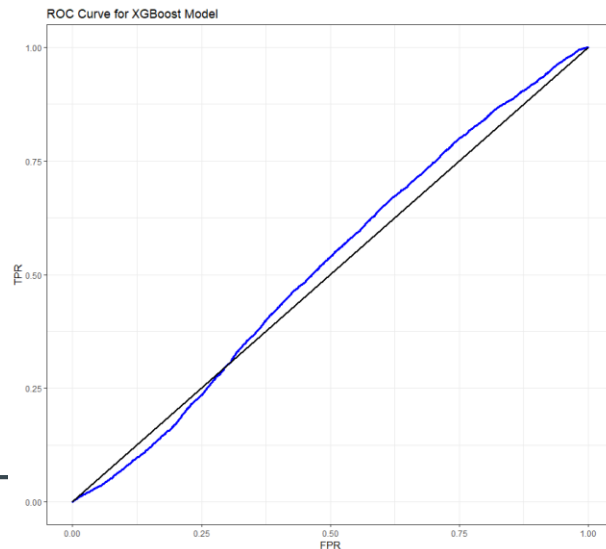
```
Residual standard error: 0.4645 on 18112 degrees of freedom
Multiple R-squared(full model): 0.2311   Adjusted R-squared: 0.1231
Multiple R-squared(proj model): 0.05514   Adjusted R-squared: -0.07752
F-statistic(full model, *iid*): 2.14 on 2543 and 18112 DF, p-value: < 2.2e-16
F-statistic(proj model): 81.44 on 15 and 5 DF, p-value: 6.074e-05
```


XGBoost

	Reference	
Prediction	0	1
0	1968	6162
1	1714	7118

Accuracy : 0.5357
95% CI : (0.5281, 0.5432)

```
> lift1
deciles y_data_test mean_response
1:      1  0.4502062    0.8646151
2:      2  0.4752358    0.9126841
3:      3  0.5247642    1.0078026
4:      4  0.5212264    1.0010084
5:      5  0.5353774    1.0281851
6:      6  0.5495283    1.0553618
7:      7  0.5854953    1.1244359
8:      8  0.5784198    1.1108475
9:      9  0.5318396    1.0213909
10:     10  0.4549204    0.8736687
```



Logistic Regression

Best	Worst	Accuracy	F1 score
0.61	0.57	0.76	0.56

```

clf = LogisticRegression(random_state=0, solver='lbfgs', multi_class='multinomial').fit(x_train, y_train)

from sklearn.metrics import f1_score

y_pred = clf.predict(x_test)

f1_score(y_test, y_pred)

0.6815907522429262

from sklearn.metrics import accuracy_score
accuracy_score(y_test, y_pred)

0.5647149006427266

from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(y_test, y_pred)
print(confusion_matrix)

[[1676 6451]
 [ 931 7901]]

```

```

from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))

              precision    recall  f1-score   support

    0         0.64      0.21      0.31      8127
    1         0.55      0.89      0.68      8832

 accuracy          0.56   16959
 macro avg         0.60   16959
 weighted avg      0.59   16959


```

TN 1696	FP 6451
FN 931	TP 7901

```

library(dplyr)
phat <- predict(out, type = "response")
phat <- predict(out, data = ww2, type = 'response')
deciles <- ntile(phat, n = 10)
df <- data.table(deciles=deciles, phat=phat, signs=ww1$signs)
lift <- aggregate(df, by=list(deciles), FUN="mean", data=df)
lift <- lift[,c(2,4)]
lift[,3] <- lift[,2]/mean(ww1$signs)
names(lift) <- c("decile", "Mean Response", "Lift Factor")
lift

```



```
> lift
```

decile	Mean Response	Lift Factor
1	0.3296225	0.6059885
2	0.4574056	0.8409092
3	0.5104116	0.9383572
4	0.5488867	1.0090911
5	0.5840194	1.0736800
6	0.5934172	1.0909574
7	0.6305085	1.1591471
8	0.6413359	1.1790526
9	0.6852300	1.2597490
10	0.7495644	1.2941801

Decision Trees

```
from sklearn.tree import DecisionTreeClassifier

tree_clf2 = DecisionTreeClassifier(max_depth=2, random_state=42)
tree_clf2.fit(x_test, y_test)

DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=2,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0, presort=False,
                        random_state=0, splitter='best')

y_pred2 = tree_clf2.predict(x_test)
accuracy_score(y_test, y_pred2)

0.5150126776342945

f1_score(y_test, y_pred2)

0.47946949154095

cross_val_score(tree_clf2, x_test, y_test, cv=10)

array([0.40518542, 0.408132, 0.59316038, 0.4054245, 0.4120283,
       0.40495282, 0.41497442, 0.40117994, 0.40530973, 0.420059 ])

tree_clf10 = DecisionTreeClassifier(max_depth=10, random_state=42)
tree_clf10.fit(x_test, y_test)


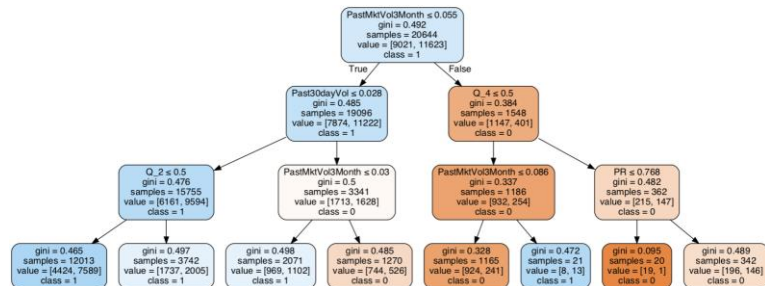
DecisionTreeClassifier(class_weight=None, criterion='gini', max_depth=10,
                        max_features=None, max_leaf_nodes=None,
                        min_impurity_decrease=0, min_impurity_split=None,
                        min_samples_leaf=1, min_samples_split=2,
                        min_weight_fraction_leaf=0, presort=False,
                        random_state=0, splitter='best')

y_pred10 = tree_clf10.predict(x_test)
accuracy_score(y_test, y_pred10)

0.7128958075358217
```

```
acc, fl, cross_min, cross_max = [], [], [], []
for i in np.arange(3,20,1):
    tree_clf5 = DecisionTreeClassifier(random_state=42,criterion='gini', max_depth = i)
    tree_clf5.fit(x_train,y_train)
    y_pred5 = tree_clf5.predict(x_test)
    print(i, accuracy_score(y_test, y_pred5), f1_score(y_test, y_pred5), \
          min(cross_val_score(tree_clf5, x_test, y_test, cv=10)), \
          max(cross_val_score(tree_clf5, x_test, y_test, cv=10)))
    acc.append(round(accuracy_score(y_test, y_pred5),4))
    fl.append(round(f1_score(y_test, y_pred5),4))
    cross_min.append(round(min(cross_val_score(tree_clf5, x, y, cv=10)),4))
    cross_max.append(round(max(cross_val_score(tree_clf5, x, y, cv=10)),4))

3 0.5461406922577983 0.46830031712038219 0.5990566037735849 0.671386430678466
4 0.5725573441830296 0.4204115829711473 0.6267688679245284 0.671386430678466
5 0.5537472728344832 0.443823738450604 0.6206489675516225 0.6821933962264151
6 0.51211714597558819 0.40996099009090909 0.6281673541543901 0.6814037735849056
7 0.52184680700513 0.4183820415078358 0.629345904537419 0.6892688679245284
8 0.5211981838551801 0.59553695953696 0.6202830188679245 0.6892688679245284
9 0.5169524505100537 0.4020596521908093 0.624042427813789 0.6827330188679245
10 0.5357037561176956 0.5810365010109609 0.6140247495580436 0.679425830188679
11 0.522731293118698 0.5597258485639687 0.6114386792452831 0.679386792452831
12 0.5163040273601038 0.5280478683620046 0.609669811207547 0.668620754716981
13 0.5056239754702518 0.5306324859941015 0.5955188679245284 0.655070754716981
14 0.5119995282740728 0.546470846102228 0.598466981120755 0.6533018867924528
15 0.5053363995518604 0.52591127473712914 0.5880866411314084 0.637971698112075
16 0.502977769915679 0.517488121815788 0.5934001718550383 0.6344339622641509
17 0.5089922754879415 0.5380817662395296 0.5875073659398939 0.6344339622641509
18 0.5068695088153783 0.5316157938952674 0.5849056603773585 0.6344339622641509
19 0.5089922754879415 0.5325849003648611 0.5837264150943396 0.637820754716981
```



Cross validation = 10,
worse results longer depths
-> overfitting. 4, 5 are the
'best' in out-of-sample, but
overall the same in cross-
validation, pick priority

Max depth	Best	Worst	Accuracy	F1 score
3 gini	0.67	0.60	0.546	0.682
5 gini	0.62	0.68	0.55	0.66
3 entropy	0.67	0.59	0.543	0.683

Random Forest Classification

Max depth	Best	Worst	Accuracy	F1 score
3 gini	0.65	0.60	0.553	0.688
10 gini	0.67	0.65	0.559	0.662
3 entropy	0.67	0.59	0.541	0.550

```
from sklearn.ensemble import RandomForestClassifier
rnd_clf = RandomForestClassifier(random_state=42,max_depth=3)

rnd_clf.fit(x_train, y_train)
y_pred = rnd_clf.predict(x_test)
accuracy_score(y_test, y_pred)

/Users/ycli/anaconda3/lib/python3.7/site-packages/sklearn/ensemble/for
n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)

0.5525479580163925

f1_score(y_test, y_pred)

0.6880190773785051

cross_val_score(rnd_clf, x_test, y_test, cv=10)

array([0.6010607, 0.63523866, 0.62971698, 0.60318396, 0.61674528,
       0.61851415, 0.61025943, 0.65663717, 0.63362822, 0.64070796])

rnd_clf2 = RandomForestClassifier(random_state=42,max_depth=10)
rnd_clf2.fit(x_train, y_train)
y_pred2 = rnd_clf2.predict(x_test)
accuracy_score(y_test, y_pred2)

/Users/ycli/anaconda3/lib/python3.7/site-packages/sklearn/ensemble/for
n_estimators will change from 10 in version 0.20 to 100 in 0.22.
"10 in version 0.20 to 100 in 0.22.", FutureWarning)

0.5599976413703638

cross_val_score(rnd_clf2, x_test, y_test, cv=10)

array([0.62816735, 0.67000889, 0.64109863, 0.64917453, 0.6692217,
       0.63860889, 0.64622642, 0.6460767, 0.6580492, 0.6259587 ])


f1_score(y_test, y_pred2)

0.6622612473974835
```

Best performance:
max_depth 8.9 gini,
greater is overfitting

```
import warnings
warnings.filterwarnings('ignore')
acc, f1, cross_min, cross_max = [], [], [], []
for i in np.arange(3,20,1):
    rnd_clf = RandomForestClassifier(random_state=42,criterion='entropy', max_depth = i)
    rnd_clf.fit(x_train, y_train)
    y_pred = rnd_clf.predict(x_test)
    print(i, accuracy_score(y_test, y_pred), f1_score(y_test, y_pred), \
          min(cross_val_score(rnd_clf, x_test, y_test, cv=10)), \
          max(cross_val_score(rnd_clf, x_test, y_test, cv=10)))
    acc.append(round(accuracy_score(y_test, y_pred),4))
    f1.append(round(f1_score(y_test, y_pred),4))
    cross_min.append(round(min(cross_val_score(rnd_clf, x, y, cv=10)),4))
    cross_max.append(round(max(cross_val_score(rnd_clf, x, y, cv=10)),4))
```

```
3 0.5523910607936788 0.6881137269403016 0.5992928697701827 0.6513274336283186
4 0.5548676219116693 0.6807898854074168 0.6096698113207547 0.6631268436578172
5 0.5546317589480512 0.681697509482068 0.6205067766647024 0.6647024160282852
6 0.5540421015390058 0.6782604841111116 0.628756629459045 0.67118443134944
7 0.5657173182381037 0.6738697250143912 0.628756629459045 0.6702064896755162
8 0.5599386756294593 0.6805632838248513 0.6340601060695344 0.6745283018867925
9 0.5645969691609175 0.673013904879993 0.6317030053034767 0.6721698113207547
10 0.5644200719182039 0.64610379479649429 0.6299351797289334 0.6774764150943396
11 0.5489120820803114 0.6555915721231766 0.6216853270477313 0.6709905660377359
12 0.555929005247951 0.6527092460225963 0.6216853270477313 0.6851415094339622
13 0.559464949702223 0.6480425872709286 0.6057748968788415 0.6682380671773719
14 0.551031048487528746 0.6378789142945951 0.6157962571525849 0.6686320754716983
15 0.5558700395070464 0.6449849170437406 0.6311137301119623 0.6586084905660378
16 0.550445191348292 0.6256138283244943 0.5969357690041249 0.6627358490566038
17 0.5444306857715667 0.6217936166046603 0.6205067766647024 0.6739386792452831
18 0.5454648964305619 0.6194122005433439 0.6149744150943396 0.6462735849056604
19 0.5418361931717672 0.6146980065456709 0.6031839622441509 0.660969811320755
```



Bagging

```
from sklearn.model_selection import cross_val_score
from sklearn.ensemble import BaggingClassifier, ExtraTreesClassifier, RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score

seed = 42
# Create classifiers
rf = RandomForestClassifier()
et = ExtraTreesClassifier()
knn = KNeighborsClassifier()
lg = LogisticRegression()
clf_array = [rf, et, knn, lg]

for clf in clf_array:
    vanilla_scores = cross_val_score(clf, x_train, y_train, cv=10, n_jobs=-1)
    clf.fit(x_train, y_train)
    y_pred = clf.predict(x_test)
    fscore=f1_score(y_test,y_pred)
    bagging_clf = BaggingClassifier(clf, max_samples=0.4, max_features=10, random_state=seed)
    bagging_scores = cross_val_score(bagging_clf, x_train, y_train, cv=10, n_jobs=-1)
    bagging_clf.fit(x_train,y_train)
    y_pred2 = bagging_clf.predict(x_test)
    fscore2=f1_score(y_test,y_pred2)

    print ("Mean: {1:.3f}, std: (+/-) {2:.3f} [{0}]" .format(clf.__class__.__name__,\
        vanilla_scores.mean(), vanilla_scores.std()))
    print ("Mean: {1:.3f}, std: (+/-) {2:.3f} [Bagging {0}]" .format(clf.__class__.__name__,\
        bagging_scores.mean(), bagging_scores.std()))
    print ("f1_score: {1:.3f} [{0}]" .format(clf.__class__.__name__, fscore))
    print ("f1_score: {1:.3f} [Bagging {0}]\n" .format(clf.__class__.__name__, fscore2))

Mean: 0.590, std: (+/-) 0.020 [RandomForestClassifier]
Mean: 0.621, std: (+/-) 0.017 [Bagging RandomForestClassifier]
f1_score: 0.546 [RandomForestClassifier]
f1_score: 0.651 [Bagging RandomForestClassifier]

Mean: 0.585, std: (+/-) 0.019 [ExtraTreesClassifier]
Mean: 0.625, std: (+/-) 0.017 [Bagging ExtraTreesClassifier]
f1_score: 0.544 [ExtraTreesClassifier]
f1_score: 0.637 [Bagging ExtraTreesClassifier]

Mean: 0.525, std: (+/-) 0.010 [KNeighborsClassifier]
Mean: 0.550, std: (+/-) 0.008 [Bagging KNeighborsClassifier]
f1_score: 0.570 [KNeighborsClassifier]
f1_score: 0.615 [Bagging KNeighborsClassifier]

Mean: 0.592, std: (+/-) 0.014 [LogisticRegression]
Mean: 0.584, std: (+/-) 0.010 [Bagging LogisticRegression]
f1_score: 0.683 [LogisticRegression]
f1_score: 0.681 [Bagging LogisticRegression]
```

Voting - soft/hard

```
# hard/soft voting

import warnings
warnings.filterwarnings("ignore")

from sklearn.ensemble import VotingClassifier
eclf = VotingClassifier(estimators=[('Random Forests', rf), ('Extra Trees', et), \
                                   ('KNeighbors', knn), ('Logistic Classifier', lg)], voting='hard')
eclf2 = VotingClassifier(estimators=[('Random Forests', rf), ('Extra Trees', et), \
                                   ('KNeighbors', knn), ('Logistic Classifier', lg)], voting='soft')
for clf, label in zip([eclf, eclf2], ['Voting Hard', 'Voting Soft']):
    scores = cross_val_score(clf, x_train, y_train, cv=10, scoring='accuracy')
    clf.fit(x_train, y_train)
    y_pred = clf.predict(x_test)
    fscore = f1_score(y_test, y_pred)

    print("Mean: {1:.3f}, std: (+/-) {2:.3f} [{0}]" .format(label, scores.mean(), scores.std()))
    print ("f1_score: {1:.3f} [{0}]\n".format(label, fscore))
```

Mean: 0.595, std: (+/-) 0.020 [Voting Hard]
f1_score: 0.559 [Voting Hard]


Mean: 0.602, std: (+/-) 0.017 [Voting Soft]
f1_score: 0.635 [Voting Soft]

Decision Boundaries

```
import matplotlib.pyplot as plt
from mlxtend.plotting import plot_decision_regions
import matplotlib.gridspec as gridspec
import itertools

gs = gridspec.GridSpec(3, 3)
fig = plt.figure(figsize=(14, 12))
labels = ['Random Forest', 'Extra Trees', 'KNN', 'Logistic', 'Voting Hard', 'Voting Soft']
for clf, lab, grd in zip([rf, et, knn, lg, eclf, eclf2], labels, \
                        itertools.product([0, 1, 2], repeat = 2)):
    clf.fit(x_train[['Ret5DayBefore', 'Past200SK']], y_train)
    ax = plt.subplot(gs[grd[0], grd[1]])
    fig = plot_decision_regions(X=np.array(x_train[['Ret5DayBefore', 'Past200SK']]), \
                              y=np.array(y_train), clf=clf)
    plt.title(lab)
```

Decision Boundaries



References

- Ball, R., Kothari, S., 1991. Security returns around earnings announcements. *The Accounting Review* 66, 718–738.
- Barber, B., De George, E., Lehavy, R., Trueman, B., 2013. The earnings announcement premium around the globe. *Journal of Financial Economics* 108, 118–138.
- Frazzini, A., Lamont, O., 2007. The earnings announcement premium and trading volume. NBER Working Paper No. 13090
- McNichols, M., 1988. A comparison of the skewness of stock return distributions at earnings and non-earnings announcement dates. *Journal of Accounting and Economics* 10, 239–273.